

CAR DHEKO-USED CAR PRICE PREDICTION

(REGRESSION MODEL)

PROJECT REPORT

BY

M.KARAN KUMAR

Introduction:

In this project, the task is to enhance the customer experience and optimize the pricing process for used cars by developing a machine learning model. This model, based on historical car price data, will take into account various features such as the car's make, model, year, fuel type, and transmission type. The goal is to predict the prices of used cars accurately and integrate this model into an interactive web application using Streamlit. The app will allow customers and sales representatives to input car details and receive real-time price predictions, making the car buying and selling process more efficient.

A. Approach:

1. Import Files and Data Wrangling:

- Load datasets from multiple cities, which are in an unstructured format (likely messy, containing text, JSON-like entries, or other formats) from Excel files.
- Use libraries like pandas to import data from each city's dataset for further processing.
- For unstructured data that includes JSON-like structures (e.g., dictionaries or lists embedded as strings), use `ast.literal_eval` to safely evaluate these strings into Python objects.
- Once JSON-like data is parsed, use `pandas.json_normalize()` to convert nested JSON objects into a flat, structured DataFrame.
- This ensures that all nested values are properly spread across multiple columns, creating a clean and organized structure.
- After converting each city's dataset into a structured format, add a new column named 'City' and assign the respective city name as the value for all rows in the dataset.
- Use `pandas.concat()` to merge the structured datasets from all cities into a single DataFrame.
- Ensure that columns are aligned and duplicate entries are handled to create a unified dataset for model training, then save the DataFrame to CSV.

2. Handling Missing Values and Data Cleaning:

- Use `pandas.dropna()` to remove rows or columns containing missing data.
- Remove symbols and units (₹, Lakh, Crore, kmpl, CC); eliminate commas and extra whitespace.
- Convert values to a numerical format suitable for machine learning models.

3. Data Visualization:

- Exploratory Data Analysis (EDA) was performed to examine the relationships between various features and the target variable (price). This analysis helped uncover key patterns and identify potential outliers.
- **Correlation Matrix:** A heatmap of the correlation matrix highlighted significant correlations between dependent and independent features, such as model year and kilometers with price.
- **Outlier Detection:** Outliers in the price column were detected using the Interquartile Range (IQR) method to prevent them from affecting the model's performance.

4. Feature Selection:

- **Categorical Features:** Attributes such as fuel type, body type, brand, insurance validity, color, city, and transmission.
- **Numerical Features:** Variables like owner number, model year, kilometers driven, mileage, engine, and seats were cleaned and converted to integers.

5. Encoding and Scaling:

- **One-Hot Encoding** is used to convert categorical data into numerical form, which is often required by machine learning models.
- **Standard Scaler** ensures that all numerical features are on the same scale, preventing certain features from dominating others.

B. Model Development

1. Train-Test Split:

- Split the dataset into training and testing sets to evaluate model performance.
- Common split ratios are 70-30 or 80-20.

2. Model Selection:

i) Linear Regression:

- **Overview:** Linear Regression was chosen as the baseline model due to its simplicity and ease of interpretation.
- **Regularization:** Ridge and Lasso regression were applied to prevent overfitting.

ii) Gradient Boosting Regressor (GBR):

- **Overview:** GradientBoostingRegressor is a machine learning technique used for regression tasks. It is part of the ensemble learning family.
- Gradient Boosting is based on the idea of fitting new models to the residuals (errors) of the previous models.

iii) Decision Tree Regressor:

- Decision Trees were chosen for their interpretability and their ability to model nonlinear relationships.
- Pruning was applied to prevent overfitting by limiting the tree depth.

iv) Random Forest Regressor:

- A Random Forest is an ensemble of decision trees. In the case of regression, each decision tree predicts a value, and the final prediction is the average of all the trees' outputs.
- Each tree is trained on a different random subset of the data (using bootstrapping).
- For each node in a tree, only a random subset of features is considered for splitting.

3. Model Evaluation:

- **Mean Squared Error (MSE):** Measures the average squared difference between actual and predicted values.
- **Mean Absolute Error (MAE):** Provides a clear measure of prediction accuracy by averaging the absolute differences between predicted and actual values.
- **R² Score:** Indicates how well the independent variables explain the variance in the dependent variable.

Model Comparison:

Model	MAE	MSE	RMSE	R2
Linear Regression	28900637467349.035	7.265332441566777e+28	269542806276976.66	5.567520170165592e+27
Decision Tree Regressor	1.0674196185286104	2.9624358446866483	1.7211728108143727	0.7729846300803596
Random Forest Regressor	0.7830135690694408	1.5521169012882599	1.2458398377352764	0.8810592326796024
Gradient Boosting Regressor	1.1081097106192699	2.48090160139723	1.5750878075197046	0.8098852648868357
Ridge Regressor	1.745035736603043	0.8494434437042098	1.320998007796773	0.8662756076292903
Lasso Regressor	1.6760714498553486	0.8446834733425146	1.2946317815716362	0.8715604319725254

4. Results:

Random Forest:

- Achieved the best performance with the highest R^2 and the lowest MSE/MAE, making it the chosen model for deployment.
- **Hyperparameter Tuning:** Grid Search was employed to identify the optimal parameters, such as `n_estimators` and `max_depth`. By systematically testing a range of values for these parameters, Grid Search helped determine the best combination that enhances the Random Forest model's performance.

5. Pipeline:

- **Modular Structure:** The pipeline is structured in a modular fashion, allowing for clear separation between data preprocessing and model training. This enhances the readability and maintainability of the code, making it easier to modify individual components without affecting the overall workflow.
- **Data Preprocessing:** The pipeline incorporates two distinct preprocessing steps for handling different types of data: numerical and categorical.
 - Numerical features are scaled using the `StandardScaler`, while categorical features are transformed using the specified encoder. This ensures that each type of data is processed appropriately to improve model performance.
- **ColumnTransformer Usage:** By utilizing `ColumnTransformer`, the pipeline efficiently applies different preprocessing techniques to specified columns in the dataset. This allows for simultaneous handling of multiple feature types, optimizing the data preparation process and ensuring that the model receives data in the correct format.
- **Integration of Model:** The final model, trained using the Random Forest regression algorithm, is integrated into the pipeline. This means that the model is directly trained on the preprocessed data, streamlining the workflow from data input to predictions and making it straightforward to apply the same preprocessing steps to new data when making predictions.

- **Reproducibility and Consistency:** By encapsulating the entire workflow (from preprocessing to model training) in a single pipeline, the approach ensures reproducibility. It guarantees that the same preprocessing steps are applied consistently during both training and inference, which is crucial for achieving reliable predictions and validating model performance.

C. Model Deployment - Streamlit

- Streamlit is an open-source Python library designed for quickly building custom web applications tailored for data science and machine learning tasks. Its ease of use and adaptability make it an excellent choice for deploying machine learning models in interactive applications.

1. Features of the Application:

i) User Input Interface:

- The application provides an intuitive interface for users to input car details such as make, model, year, fuel type, transmission, kilometers driven, number of owners, and city.
- Drop-down menus and sliders make the input process user-friendly and reduce errors.

ii) Price Prediction:

- Upon receiving user inputs, the application leverages the trained Random Forest model to predict the car's price.
- The predicted price is displayed instantly, enhancing the user experience.

Backend Implementation:

iii) Model Loading:

- The trained Random Forest model, StandardScaler, and LabelEncoder are loaded into the application using the pickle library, ensuring they are ready for predictions.
- **Data Preprocessing:** User inputs are preprocessed in the same way as the training data, ensuring consistency and accuracy in predictions.

D. Conclusion - Project Impact:

- Deploying the predictive model through the Streamlit application revolutionizes the user experience at CarDekho by delivering swift and reliable price estimates for used cars.
- This real-time pricing tool empowers customers with data-driven insights, enabling them to make informed decisions when buying or selling vehicles. For sales representatives, it simplifies the valuation process, saving time and ensuring consistency across pricing.
- Moreover, this deployment lays the groundwork for future innovations, where advanced features like personalized recommendations or integration with real-time market data can further enhance the accuracy and sophistication of the predictions.